

Reminder ask — verification round three

Re-walkthrough after PR #1608, joshing-11.vercel.app

2026-09-06 · Duo Prova (+1 555-222-2222, inviter) → Rue Prova (+1 555-333-3333, fresh signup) · Production at af3582ad · Verifies #1608; re-confirms #1607 and #1610

Third pass through invite → signup → onboarding → reminder ask against production. The confirmation toast — invisible in round one, invisible again in round two — now reaches the player, and this time the handoff from the loading screen was measured rather than eyeballed. Four further observations came out of the walkthrough, none of them in the flow it set out to test.

HEADLINE · THE FIX WORKED

The toast is visible, and it waits for the loader to get out of its way. A 50 ms-resolution DOM probe over the deliberate reproduction caught the exact handoff #1608 was written to produce: at $t=0$ the fullscreen loader is up and the toast is *not in the DOM at all* — the `!loading` gate holding it back rather than painting it under the takeover layer. At $t=1103$ ms the loader is gone and the toast is present in the same tick, at `z-index 70`, `opacity 1`, and — the part neither earlier round could claim — `elementFromPoint` at its own centre returns the toast itself, so nothing is covering it. It was also caught on screen, over the first question.

WHAT ROUND THREE CONFIRMED

BEHAVIOR	RESULT
Confirmation toast visible to the player	✓ Fixed — the defect from rounds one and two is closed
Toast held back while the loader is up, released when it clears	✓ Confirmed by probe, 50 ms resolution
Invite-link rows read “Swap”, not a bare ×	✓ Still live
At-cap helper: “All 3 topics are taken. Tap Swap above...”	✓ Still live
“No category — carries all 3 topics above”	✓ Still live, accurate count
Invite-only signup admits a fresh number through a link	✓ Confirmed — <code>joined_via_invite_link_id</code> set
Link-invite topics arrive unselected as suggestions	✓ Confirmed
SMS opt-in completes end-to-end with a test number	✓ Confirmed — <code>opted_in</code>
Opt-in retires both later prompts in one write	✓ Confirmed — identical timestamps
No Twilio send fires for an allowlisted number	✓ Confirmed — no new <code>SmsLog</code> row
Crafting screen on the live pass	· Not exercised — see the caveat under Steps

STEPS

1. Logged in as **Duo Prova**, opened the invite-link panel, and re-confirmed all three #1607 fixes render as intended. Cancelled without creating a second link.
2. Logged out and opened Duo's existing /u/duoprova/... link as a number the app had never seen since the prior fixture was deleted.
3. Signed up +1 555-333-3333; OTP bypass carried it with no real text. Profile saved as **Rue Prova**.
4. Onboarding arrived with all three of Duo's seeded topics as unselected suggestions. Selected all three; generation started here, at 15:41:04Z.
5. Reached the reminder ask and tapped **Text me when they open** at 15:45:56Z. It succeeded, straight through to /daily, **and the toast was there** — first line of the page's own text, over question one.
6. *Caveat on that pass*: the build had finished four and a half minutes earlier, because inspecting the ask screen kept me on it far longer than a player would linger. So the live pass took the fast path and the crafting screen never appeared — the same fast path round one hit, for a different reason. It does not exercise the loader handoff.
7. So the handoff was reproduced deliberately via /daily?remindersOn=1, the same route round two used to prove the failure, this time with a DOM probe armed at page load rather than a screenshot race.
8. Played 2 of 5, then verified consent, suppression state and the SMS log directly in the database.

EVIDENCE

The handoff, as recorded by the probe (entries are state changes, not samples):

```
t=0      ms    loader=true   toast=absent
t=1103   ms    loader=false  toast=present  z=70  opacity=1  onTop=true
t=4334   ms    loader=false  toast=absent   (self-dismissed)
```

The dismissal timestamp is the one number here not to trust: a screenshot capture blocked the page between the second and third entries, so the 1.8 s timer is observed late. Appearance and coverage — the two things in question — are unaffected.

The opt-in write:

sms_opt_in	opted_in
sms_opt_in_at	2026-09-06 15:45:56.617+00
sms_consent_source	onboarding_web_form
sms_consent_policy_version	2026-09-02
reminder_prompt_dismissed_at	2026-09-06 15:45:56.815+00
reminder_interstitial_seen_at	2026-09-06 15:45:56.815+00
joined_via_invite_link_id	2a0730bc... (Duo's standing link)

Every SMS ever logged for that number — still only round one's two failed real sends, none from this signup or this opt-in:

```
otp    2026-09-05 23:25:47+00
otp    2026-09-05 23:25:06+00
```

WHAT ELSE THE WALKTHROUGH TURNED UP

OVERCLAIM · ACTIVITY COPY

“You’ve got Renaissance Florence down cold” — on two correct answers, at the lowest tier.

Home said it on a first-ever session. `PLAYER_MASTERY` for that territory reads `tier=establishing`, `tier_reached_at=null`, 60 points; the session recap on the same data says `NEXT: FAMILIAR`. Two surfaces, opposite claims, same moment.

It is one of six lines in the `you_got` pool in `activity-stream.ts`, chosen by hash of the event id. The other five describe the event that happened (“You knew {friend}’s {topic} question”); this one alone asserts a standing fact about the player’s knowledge that the event does not evidence. Same shape as the “First five complete” eyebrow #1610 just fixed — copy asserting a state its gate never checks. Cheapest fix is to retire the line from the pool.

MISLABEL · QUESTION BADGE

The player’s first two questions were badged “General Knowledge”, not “Renaissance Florence”. Both slots carry domain: “Renaissance Florence” — the topic they had just chosen — and both display the generic bucket instead, on the screen whose promise is “Made from your topics, not pulled off a shelf.”

Not a data error at the slot level: `questionBadges` in `daily/page.tsx` prefers `broad_category`, then `category`, and only falls back to `domain` — least specific first. 68 of 1,302 questions with a specific `canonical_subcategory` carry “General Knowledge” as their broad category (5.2%), and every one of them badges generically. Hitting two of two on a first session was luck; the exposure is real. Either reorder the fallback, or stop letting a specific subcategory sit under a generic broad category.

LEAK SHAPE · OUTSIDE #1611’S REACH

The first question a new player saw glossed its own answer. “Literally meaning *the offices*, what famous Florentine museum...” — answer: **The Uffizi**, which is what *uffizi* means. The stem hands over a translation of the answer word while withholding the word.

It is `answer_source: creator_written`, so it never passed through the generator at all. The authored path calls `textContainsAnswer`, a containment check the stem passes cleanly; the partial-leak gate #1611 shipped lives in `generate-questions.ts` and does not run here. Worth deciding whether the new gates belong on the authored path too — the leak family is identical, and a friend-authored question is exactly what a new player meets first.

BONUS · THE BUILD-LATENCY DOC IS UNBLOCKED

This signup produced the first post-deferral built row. The diagnosis doc’s Phase 2 was waiting on data “that arrives on its own at the next cron”; the walkthrough generated it. `npm run check:build-latency` now passes every Phase 2 check: `deferred=true`, `span_ms 22282`, `user_visible_ms 21022`, target size and call attribution both clean.

Phase 3’s subtraction reads **1,260 ms saved** against a pre-registered prediction of ~7,646 ms — a sixth of it. On one row that is an instrument reading, not a population estimate, and the doc says so. But it is the first evidence on the question, and it points the other way from the prediction.

INVESTIGATED AND CLEARED · #1610 HOLDS

Home showed **“Catch up · 5 missed questions”** on a first-day account, which is exactly the contradiction #1610 shipped to remove. It is not a regression. The 17:00 UTC daily reset landed mid-session, so by the time home rendered, the round I had been playing was genuinely the previous window’s and its five unplayed slots were genuinely missed. `isCatchUpSlotEligible` behaved as written. Two boundary artifacts came with it, both from a session that straddled the reset and neither reproducible on demand: two answers landed against the expired round and were credited (they survive in `MASTERY_EVENTS` and in the recap’s “+60 this round”), and the queue row was then rebuilt in place for the new day, so the round’s own dots went back to unanswered. Noted rather than filed — one observation each, at a boundary a real player crosses rarely.

WHAT THIS SAYS ABOUT THE PROCESS

THREE ROUNDS, THREE DIFFERENT FAILURE MODES

Round one found the toast never rendered. Round two found the fix rendered it under a takeover layer. Round three found it working — and to say so with any confidence needed an instrument, not a look, because the thing being verified lives for 1.8 seconds and the previous round’s failure was invisible by construction.

The general lesson is narrow but worth keeping: when the defect is *“the player never saw it”*, the verification has to measure presence and coverage, not take a screenshot and hope to catch it.

THE FLOW UNDER TEST WAS THE LEAST OF IT

The reminder ask passed cleanly. Everything worth Josh’s attention above — a mastery claim on two answers, a topic badged as its own opposite, an answer glossed in its own stem — came from walking past those screens on the way, on a genuinely new account. None of it is reachable from a test suite, and none of it was what the walkthrough was for.